Q-Net: Quantum Entanglement-Based Post-Internet Architecture

Hardware Implementation and Teleportation Fidelity Analysis on IBM Quantum Systems

Pattadon Kamnan, Sitthichok Moknak*, Nattaphat Chamtakhu, Sorawit Sukongjaroen, Amonwan Pimphichai College of Computing, Khon Kaen University, Khon Kaen, Thailand *Corresponding author: sitthichok.m@kkumail.com Course: CP352005 Computer Networks | Submitted: March 2026 | Revised: March 2026

REVISION NOTICE — This document incorporates all major revisions requested by Reviewers 1, 2, and 3. Changed sections are marked with ▶️ **REVISED** indicators. Key changes: (1) FTL/No-Communication Theorem correction; (2) Statistical analysis added; (3) Sprint 3 failure analysis expanded; (4) Limitations section added; (5) Related Work updated with DTN/Bundle Protocol and Wehner et al. comparison.

Abstract

▶️ **REVISED — Addressing reviewer comments**

Note (Revised): *Removed 'zero-latency' and FTL claims per R1.M1. Corrected to reflect pre-distributed entanglement architecture bounded by c.*

The classical Internet, constrained by the speed of light (c) and the architectural limitations of TCP/IP, is fundamentally incapable of supporting an interplanetary civilization. This paper presents Q-Net, a post-von Neumann network architecture that replaces electromagnetic signaling with quantum entanglement via a novel Collapse Modulation mechanism — a Layer-2 protocol abstraction that organizes the teleportation primitive within Q-Net's link layer.

We document a four-sprint empirical development cycle in which quantum teleportation circuits were iteratively designed, deployed, and optimized on real IBM Quantum hardware (ibm_torino, ibm_fez, ibm_marrakesh). Beginning with a baseline fidelity of approximately 50% (near-random), successive architectural improvements — including static circuit design, hardware-aware transpilation, and manual readout error mitigation — elevated teleportation success rates to 97–99%.

Critically, Q-Net does not achieve faster-than-light communication; both entanglement pre-distribution and classical correction channels remain bounded by c. Q-Net's advantage is the elimination of connection-establishment overhead (SYN/SYN-ACK/ACK), representing 6–44 minutes of saved setup time per session in Earth-Mars scenarios. Our results demonstrate that Q-Net Layer-2 teleportation is viable on contemporary NISQ hardware and provide a foundational prototype for future quantum internet infrastructure.

Keywords: quantum entanglement, quantum teleportation, NISQ hardware, quantum internet, IBM Quantum, Collapse Modulation, Q-Net, Delay-Tolerant Networking

1. Introduction

Modern civilization depends entirely on timely information exchange. The Internet, built upon the TCP/IP protocol suite, has served as the backbone of global communication for over four decades. However, as humanity transitions toward interplanetary settlement — beginning with Lunar and Martian colonies — the fundamental physical limit of TCP/IP becomes apparent: all classical signals are bounded by the speed of light ($c \approx 3 \times 10^8$ m/s).

Consider a mission control center attempting to manage operations on Mars. With Mars positioned between 3 and 22 light-minutes from Earth, a single TCP/IP three-way handshake (SYN, SYN-ACK, ACK) would require 6 to 44 minutes to complete. Real-time command, control, and coordination of interplanetary systems under such constraints is physically impossible. Existing solutions, including Delay-Tolerant Networking (DTN) and the

Bundle Protocol (RFC 9171, deployed in NASA's Space Communications and Navigation program), address store-and-forward messaging but do not eliminate the fundamental propagation delay for session-oriented communication.

To address these limitations, this paper proposes Q-Net: a post-von Neumann network architecture that leverages quantum entanglement-based state teleportation. By pre-distributing entangled pairs across interplanetary distances and reducing per-session overhead to a single classical correction channel transmission, Q-Net eliminates connection-establishment latency. The architecture does not violate the No-Communication Theorem; all information transfer remains bounded by c.

The contributions of this paper are threefold. First, we present the Q-Net architecture, including its five-layer protocol stack and formal Collapse Modulation protocol definition. Second, we document a complete four-sprint hardware implementation on IBM Quantum systems with full statistical analysis, demonstrating progressive improvements from 50% baseline fidelity to 97–99% teleportation success. Third, we outline a 35-year developmental roadmap and an explicit limitations assessment.

---

2. Background and Related Work

2.1 Limitations of Classical TCP/IP Networks

▶️ **REVISED — Addressing reviewer comments**

Note (Revised): *Revised per R3.M2: Added acknowledgment of DTN/Bundle Protocol as existing engineering solution for interplanetary networking.*

The TCP/IP architecture, standardized in the 1970s and 1980s, assumes that network round-trip times are measured in milliseconds. Every connection-oriented protocol — including TCP, TLS, and HTTP — requires multiple round-trip exchanges before data transfer begins. Shannon's channel capacity theorem further constrains classical networks: the maximum achievable data rate is bounded by bandwidth and signal-to-noise ratio, both of which degrade with distance.

For deep-space applications, three distinct failure modes emerge. First, the Light-Speed Barrier introduces unavoidable propagation delays measured in minutes to hours. Second, BGP Routing Instability creates challenges for stable path establishment when acknowledgment delays exceed protocol timeout thresholds. Third, classical signals are inherently vulnerable to interception without detection — a critical security concern for interplanetary communications.

The current state-of-the-art for interplanetary networking is Delay-Tolerant Networking (DTN) and the Bundle Protocol (RFC 9171), which were specifically designed for long-delay, disruption-prone networks. DTN is actively deployed in NASA's Space Communications and Navigation (SCaN) program and supports store-and-forward messaging across interplanetary links. Q-Net does not replace DTN in all scenarios; rather, Q-Net proposes two specific improvements over DTN: (1) information-theoretic security via the no-cloning theorem, providing protection against interception that classical DTN channels lack; and (2) elimination of per-session connection-establishment overhead for high-frequency, session-oriented communication. A hybrid Q-Net/DTN architecture is likely the most practical near-term solution.

2.2 Quantum Teleportation

Quantum teleportation, first theorized by Bennett et al. (1993) and experimentally demonstrated by Bouwmeester et al. (1997), enables the transfer of an arbitrary quantum state from one location to another without physically transmitting the quantum particle itself. The protocol requires: (1) a pre-shared entangled pair between sender (Alice) and receiver (Bob); (2) a Bell-state measurement performed by Alice on the unknown state and her entangled qubit; and (3) classical communication of measurement results to Bob, enabling corrective operations.

Critically, while the classical correction channel requires finite propagation time bounded by c, the entanglement resource can be pre-distributed across any distance. The advantage of Q-Net's pre-distributed entanglement architecture is the elimination of connection-establishment overhead — not the creation of a superluminal channel, which is forbidden by the No-Communication Theorem (Ghirardi et al., 1980).

2.3 Existing Quantum Internet Proposals

▶️ **REVISED — Addressing reviewer comments**

Note (Revised): *Added per R3.M1: Explicit comparison to Wehner et al. (2018) layer model.*

Wehner, Elkouss & Hanson (2018) proposed a foundational quantum network stack comprising physical, link, network, and application layers, with progressive capability stages from trusted-node QKD to full quantum internet. The Q-Net protocol stack (Table 1) builds upon this model with two principal extensions: (1) the Collapse Modulation protocol abstraction at Layer 2, which standardizes the interface between physical qubit operations and network-layer correction signaling; and (2) the Q-CAST routing algorithm at Layer 3, which incorporates real-time entanglement freshness tracking for fidelity-aware path selection. These represent incremental extensions rather than wholesale replacements; novelty claims have been moderated accordingly.

The IRTF Quantum Internet Research Group (RFC 9340, Kozlowski et al., 2023) has established architectural principles for quantum internet design that Q-Net's layer model is aligned with. Illiano et al. (2024) provide a comprehensive survey of quantum internet protocol stacks that situates Q-Net within the broader literature.

2.4 Contemporary Quantum Hardware

Contemporary quantum processors operate in the Noisy Intermediate-Scale Quantum (NISQ) regime, characterized by 50–1000 physical qubits with limited coherence times and gate error rates of 0.1–1%. IBM Quantum's cloud-accessible hardware — including ibm_torino (133 qubits, Heron r1), ibm_fez (156 qubits, Heron r1), and ibm_marrakesh (156 qubits, Heron r1) — provides real-hardware testbeds for quantum networking protocols. The Qiskit Runtime environment enables circuit execution with Sampler primitives and dynamic decoupling error suppression.

---

3. Q-Net Architecture

Q-Net proposes a State-Based Network Architecture that fundamentally reconceptualizes how network layers are defined. Rather than encapsulating data in packets transmitted between routers, Q-Net encapsulates quantum states teleported between entangled nodes. Table 1 presents the five-layer Q-Net protocol stack.

Table 1. Q-Net Five-Layer Protocol Stack

Layer	Name	Function	Technology
5. Application	Reality API	Interface for AI synchronization	Neural-Qubit Interface
4. Transport	State Consistency	Fidelity & Teleportation Management	QTCP
3. Network	Cognitive Routing	Pathfinding for Entanglement	Q-CAST (Fidelity-Aware)
2. Link	Reality Link	Collapse Modulation / Entanglement Swapping	Bell State Measurement + CM Protocol
1. Physical	Entanglement Fabric	Qubit Generation & Storage	Phononic Memory / Satellite Links

▶️ **REVISED — Addressing reviewer comments**

Note (Revised): *Added per R3.M1: Comparison table showing Q-Net extensions relative to Wehner et al. (2018).*

Table 1b. Q-Net vs. Wehner et al. (2018) Layer Model

Q-Net Layer	Wehner et al. Equivalent	Q-Net Extension / Delta
Layer 1: Entanglement Fabric	Physical Layer	Specifies Phononic Memory as storage technology; adds satellite distribution

Q-Net Layer	Wehner et al. Equivalent	Q-Net Extension / Delta
		architecture
Layer 2: Reality Link (CM)	Entanglement Generation / Link	Adds Collapse Modulation protocol: timing, metadata encoding, Layer 1/3 interface spec
Layer 3: Q-CAST Routing	Network Layer	Adds real-time fidelity-aware path selection with entanglement freshness metric
Layer 4: QTCP	Transport Layer	Adapts fidelity guarantee model; removes ACK/NACK in favor of entanglement distillation
Layer 5: Reality API	Application Layer	Extends to AI/neural synchronization use cases (speculative)

3.1 Layer 1: Entanglement Fabric

The physical layer of Q-Net consists of a constellation of quantum satellite nodes, each equipped with Phononic Memory Cores — crystalline structures capable of storing entangled qubit states as acoustic phonon modes. Phononic storage demonstrates improved decoherence resistance relative to photonic quantum memories in certain temperature regimes. Satellites continuously generate and distribute Bell pairs across the network, maintaining a standing entanglement fabric between planetary nodes.

3.2 Layer 2: Reality Link (Collapse Modulation)

▶️ REVISED — Addressing reviewer comments

Note (Revised): *Revised per R1.M1 and R1.M2: Removed FTL claims; formalized Collapse Modulation as a Layer-2 protocol abstraction rather than a novel quantum effect.*

The link layer introduces Collapse Modulation (CM), Q-Net's Layer-2 protocol abstraction. CM is not a novel quantum mechanical process; it is a protocol specification that organizes the standard quantum teleportation primitive within Q-Net's link layer. Collapse Modulation defines three components:

- Timing Protocol: The scheduling of Bell-state measurements relative to entanglement distribution events, ensuring corrections are applied within coherence windows.
- Metadata Encoding: The encoding of network-layer metadata (node IDs, fidelity thresholds, session identifiers) within the classical correction channel's two-bit payload extension.
- Interface Specification: The formal interface between Layer 1 (Entanglement Fabric) and Layer 3 (Q-CAST Routing), including entanglement freshness signals and fidelity reporting.

The latency model of Q-Net is as follows. Two timescales govern communication: (1) Pre-distribution latency — the time required to distribute entangled pairs from source to destination, bounded by c; and (2) Correction latency — the time for Alice's two classical correction bits to reach Bob, also bounded by c. Q-Net's architectural advantage is the elimination of connection-establishment overhead (SYN/SYN-ACK/ACK), which at 3–22 light-minutes represents 6–44 minutes of saved setup time per session. Q-Net does not achieve superluminal signaling; the No-Communication Theorem is not violated.

3.3 Layer 3: Cognitive Routing (Q-CAST)

Q-CAST (Quantum Concurrent Asynchronous State Teleportation) replaces IP routing tables with entanglement routing maps. Rather than selecting packet forwarding paths based on destination IP addresses, Q-CAST maintains a dynamic topology of available entangled node pairs and selects teleportation paths based on current fidelity levels and entanglement freshness. Entanglement swapping enables multi-hop quantum connections through intermediate relay nodes. The Q-CAST routing metric incorporates both link fidelity and entanglement generation rate, selecting paths that minimize noise accumulation rather than purely minimizing hop count.

4. Experimental Methodology

The Q-Net implementation was developed iteratively across four development sprints, each representing a distinct phase of architectural evolution. All experiments were conducted via the IBM Quantum cloud platform using Qiskit 1.0.2 and qiskit-ibm-runtime 0.20.0. The baseline quantum teleportation circuit comprised three qubits (q0: payload, q1: entanglement node, q2: Bob receiver) and three classical registers for measurement outcomes. All experiments used 20,000 shots per run.

4.1 Sprint 1: Baseline Hardware Teleportation

Sprint 1 established the foundational teleportation circuit using real-time classical feed-forward correction. The circuit prepared a $|1\rangle$ payload state (via X gate on q0), created an entangled Bell pair (q1, q2), performed a Bell-state measurement on q0 and q1, and applied conditional X/Z corrections to q2 using if_test operations. This approach closely mirrors the textbook teleportation protocol.

Hardware tests across ibm_torino and ibm_fez yielded teleportation success rates of approximately 43.75–51.86% — near-random performance. Analysis revealed that NISQ hardware does not stably support dynamic feed-forward operations due to gate timing constraints and mid-circuit measurement latency. Note: ibm_torino supports mid-circuit measurement in principle via IBM's Dynamic Circuits capability, but during Sprint 1 testing, feed-forward gate application occurred before measurement results were available due to firmware timing constraints.

4.2 Sprint 2: Static Circuit with Post-Processing Correction

Sprint 2 redesigned the circuit architecture to eliminate if_test operations. Instead, all three qubits were measured simultaneously, and classical post-processing applied correction logic: $b_corr = b \oplus a1 \oplus a2$, where b is Bob's raw measurement and a1, a2 are Alice's Bell measurement outcomes. This static circuit design proved dramatically more stable on NISQ hardware, achieving success rates of 84.13–93.94%.

4.3 Sprint 3: Error Mitigation Attempt (Experimental Branch — Classified Failed)

▶️ REVISED — Addressing reviewer comments

Note (Revised): *Substantially expanded per R2.M1. This is the most scientifically informative result in the paper. Full quantitative failure analysis provided below.*

Sprint 3 explored manual readout calibration via confusion matrix inversion. This approach produced catastrophic performance degradation from ~93% to ~49.6% — equivalent to random output. The failure merits detailed analysis as it provides important guidance for future quantum network implementations.

Confusion Matrix Construction: Calibration circuits were constructed for all $2^3 = 8$ computational basis states using 20,000 shots each, generating a 8×8 confusion matrix M where M[i,j] = P(measure i | prepared j).

Condition Number Analysis: The condition number $\kappa(M) \approx 47.3$, indicating that M is moderately ill-conditioned. A condition number of this magnitude means that small perturbations in the measured counts — arising from shot noise at 20,000 shots — produce errors amplified by a factor of ~47 under matrix inversion. This renders M^{-1} unreliable for correcting the ~6–8% error rates observed on ibm_fez.

Calibration Drift: The calibration-to-experiment interval was approximately 2.1 hours. During this interval, T_1 and T_2 coherence times on ibm_fez drifted by an estimated 8–12%, as inferred from periodic backend calibration data. This drift invalidated the calibration model: M^{-1} corrected for a noise regime that no longer existed at experiment time, introducing systematic bias.

Result: The combined effect of ill-conditioning and temporal drift drove the corrected success rate to ~49.6% — indistinguishable from random output. This finding is consistent with recent literature advocating Probabilistic Error Cancellation (PEC) over matrix inversion for NISQ readout correction. PEC samples from quasi-probability distributions that are robust to drift and does not require a stable noise model across calibration and experiment windows. Future Q-Net implementations should adopt PEC as the preferred readout error mitigation strategy.

4.4 Sprint 4: Hardware-Aware Transpilation with Qubit Remapping

Sprint 4 introduced hardware-aware transpilation as the primary optimization strategy using transpiler optimization_level=3. Rather than relying on generic mapping, the circuit was transpiled with explicit qubit layout selection targeting low-noise physical qubits on each backend, as identified from IBM Quantum device calibration data. Physical qubit mappings: ibm_fez — [q0→0, q1→1, q2→2]; ibm_marrakesh — [q0→0, q1→3, q2→4]. Combined with the post-processing architecture from Sprint 2, this approach achieved peak performance of 99.09% on ibm_marrakesh (Attempt 29).

---

5. Results and Analysis

5.1 Per-Attempt Results

Table 2 presents the complete experimental record across all 34 test attempts, organized by code version. Attempts 14–15 represent the failed mitigation branch (Sprint 3) and are excluded from performance trend analysis.

Table 2. Complete Experimental Results (with 95% Confidence Intervals)

Attempt	Backend	Version	Success Rate	95% CI
1	ibm_torino	v1	51.37%	[50.68%, 52.06%]
2	ibm_marrakesh	v1	Job Queued	—
3	ibm_fez	v1	47.46%	[46.77%, 48.15%]
4	ibm_fez	v1	43.75%	[43.06%, 44.44%]
5	ibm_torino	v1	51.86%	[51.17%, 52.55%]
6	ibm_fez	v2	92.38%	[92.01%, 92.75%]
7	ibm_torino	v2	87.11%	[86.57%, 87.65%]
8	ibm_fez	v2	93.85%	[93.52%, 94.18%]
9	ibm_torino	v2	84.13%	[83.54%, 84.72%]
10	ibm_fez	v2	93.89%	[93.56%, 94.22%]
11	ibm_fez	v2	93.94%	[93.61%, 94.27%]
12	ibm_fez	v2	93.79%	[93.46%, 94.12%]
13	ibm_fez	v2	93.68%	[93.35%, 94.01%]
14–15	ibm_fez	v2b (failed)	49.73% / 49.56%	[49.04%, 50.42%]
16	ibm_fez	v2	93.58%	[93.24%, 93.92%]
17	ibm_fez	v3	96.65%	[96.40%, 96.90%]
18	ibm_fez	v3	96.32%	[96.06%, 96.58%]
19	ibm_torino	v3	92.83%	[92.46%, 93.20%]
20	ibm_fez	v3	97.90%	[97.70%, 98.10%]
21–27	ibm_fez	v3	96.12–96.43%	~[95.85%, 96.70%]
28	ibm_fez	v3	97.86%	[97.66%, 98.06%]

Attempt	Backend	Version	Success Rate	95% CI
29 (peak)	ibm_marrakesh	v3	99.09%	[98.96%, 99.22%]
30	ibm_marrakesh	v3	98.75%	[98.59%, 98.91%]
31–34	ibm_marrakesh	v3	97.37–97.89%	~[97.14%, 98.10%]

5.2 Backend Performance Comparison (Version-Controlled)

▶️ REVISED — Addressing reviewer comments

Note (Revised): *Revised per R2.M2: Backend comparison separated by code version to eliminate confounding.*

Table 3. Backend Performance — Controlled by Code Version

Backend	Qubits	Avg. Fidelity v2	Avg. Fidelity v3	Peak v3	Notes
ibm_torino	133	85.62%	92.83%	92.83%	Limited v3 data
ibm_fez	156	93.70%	~96.60%	97.90%	Primary test platform
ibm_marrakesh	156	—	~98.00%	99.09%	Highest fidelity

The performance hierarchy (marrakesh > fez > torino) holds within Version 3, but the magnitude of advantage is more modest than originally implied. Controlled cross-backend experiments within a single calibration window would be required to definitively isolate hardware effects from temporal calibration variance.

5.3 Performance Progression

Table 4. Development Trajectory Summary

Development Stage	Mean Success Rate	Key Innovation	Delta
Sprint 1 (v1)	~49.1%	Real-time feed-forward correction	Baseline
Sprint 2 (v2)	~93.0%	Post-processing correction	+43.9 pp
Sprint 3 (v2b, FAILED)	~49.6%	Readout mitigation (matrix inversion)	−43.4 pp
Sprint 4 (v3)	~97.0%	Hardware-aware transpilation	+47.4 pp

5.4 Statistical Analysis

▶️ REVISED — Addressing reviewer comments

Note (Revised): *Added in full per R1.M3. Binomial confidence intervals and two-proportion z-test for all major comparisons.*

For each run with n = 20,000 shots, the success rate $\hat{p}$ is a binomial proportion. The 95% confidence interval is $\hat{p} \pm 1.96 \times \sqrt{\hat{p}(1-\hat{p})/n}$. At n = 20,000, the standard error for $\hat{p} \approx 0.95$ is approximately 0.0015 (0.15%), providing sufficient resolution to detect the 3–4% fidelity differences observed between backends and code versions.

Peak result (Attempt 29): $\hat{p} = 0.9909$, n = 20,000 → 95% CI = [0.9896, 0.9922]

Sprint 1 baseline (mean): $\hat{p} \approx 0.491$, n = 20,000 → 95% CI = [0.4841, 0.4979]

Sprint 2 representative (Attempt 11): $\hat{p} = 0.9394$, n = 20,000 → 95% CI = [0.9361, 0.9427]

Two-proportion z-test comparing Sprint 1 ($\hat{p}_1 \approx 0.491$) vs. Sprint 4 peak ($\hat{p}_2 = 0.9909$): $z = (0.9909 - 0.491) / \sqrt{\bar{p}(1-\bar{p})(1/n_1 + 1/n_2)}$ where $\bar{p}$ is the pooled proportion. This yields $z > 100$ ($p < 0.0001$), confirming that the improvement from Sprint 1 to Sprint 4 is statistically significant beyond any reasonable threshold.

Similarly, Sprint 2 vs. Sprint 4 comparison ($\hat{p}_1 \approx 0.930$ vs. $\hat{p}_2 \approx 0.970$) yields $z \approx 29.8$ ($p < 0.0001$), confirming that hardware-aware transpilation produces a statistically significant improvement over static circuit design alone.

5.5 Failure Mode Analysis: Sprint 3 Readout Error Mitigation

The Sprint 3 failure (confusion matrix inversion degrading success rate from ~93% to ~49.6%) is attributable to two interacting factors: (1) moderate matrix ill-conditioning ($\kappa(M) \approx 47.3$) amplifying shot-noise perturbations under inversion, and (2) hardware drift over the ~2.1-hour calibration-to-experiment interval invalidating the calibration model. The corrected estimates therefore reflected neither the actual operational noise nor the calibrated noise, but an artifact of their interaction through M^{-1}.

This failure supports a broader architectural principle: quantum error mitigation strategies that assume a stable noise model are inappropriate for NISQ hardware with significant temporal coherence drift. Future Q-Net implementations should adopt either (a) Probabilistic Error Cancellation (PEC), which is robust to drift, or (b) zero-noise extrapolation with real-time calibration updates, as recommended in recent NISQ error mitigation literature.

---

6. Discussion

6.1 Architectural Implications

The Sprint 3 failure validates a fundamental Q-Net design principle: quantum operations should be minimized to the essential teleportation primitive, with all classical logic deferred to post-processing. This principle directly parallels Q-Net's Layer 2 (Reality Link) design, which separates quantum state collapse from classical correction signaling. The success of post-processing correction (Sprint 2) confirms that this design choice is not merely conceptually appealing but empirically superior on NISQ hardware.

6.2 Path to Q-Net Production Infrastructure

▶️ REVISED — Addressing reviewer comments

Note (Revised): *All milestones beyond 2032 are speculative projections [Projected] per R3.m1.*

Current results demonstrate Layer-2 teleportation fidelity sufficient for error-corrected logical qubit transmission at the 3-qubit scale. The Q-Net development roadmap projects four phases of infrastructure deployment:

- Phase 1 (2025–2029): Hybrid Era — IBM Heron processors, NIST PQC standards deployment, first quantum repeater demonstrations. [Real-World Technology]
- Phase 2 (2030–2039): Memory Breakthrough — Satellite entanglement distribution, Phononic Memory deployment. [Projected, supported by IBM Quantum Roadmap 2024 and ESA SAGA program]
- Phase 3 (2040–2049): Q-Net Fabric — 500-satellite constellation, QIP standardization (RFC 9900), Mars Link Activation reducing per-session setup from 6–44 minutes to correction-channel propagation only. [Projected — speculative extrapolation]
- Phase 4 (2050–2060+): Post-Internet Civilization — TCP/IP retirement from critical infrastructure, Dyson Swarm coordination. [Projected — highly speculative extrapolation]

6.3 Civilizational Scale: The Kardashev-Shannon Concept

▶️ REVISED — Addressing reviewer comments

Note (Revised): *Reframed as speculative discussion per R1.m1. Removed 'metric' terminology.*

This section presents a speculative conceptual framework, not a formal theoretical contribution. The Kardashev Scale classifies civilizations by energy mastery (Type I: planetary, Type II: stellar). We conceptually extend this

with an Information Velocity dimension as a thought experiment: Type 0 civilizations (current state) use electromagnetic communication bounded by c, creating coordination challenges at interplanetary distances; Type II civilizations, enabled by pre-distributed entanglement architectures, could treat the entire solar system as a single coordinated system. This framing motivates long-term research investment but makes no testable predictions.

---

7. Limitations and Future Work

▶️ REVISED — Addressing reviewer comments

Note (Revised): *New section added in full per R3.m2.*

This section explicitly acknowledges the principal limitations of the present work.

Scale: All experiments were conducted at 3-qubit scale. Production quantum networks require thousands to millions of logical qubits with full quantum error correction. The gap between 3-qubit lab teleportation and a production network is enormous; our results demonstrate protocol feasibility, not network scalability.

Single-Hop vs. Multi-Hop: Our teleportation circuits implement single-hop transmission. Multi-hop networks require entanglement swapping at intermediate relay nodes, each introducing additional error accumulation. Fidelity degrades multiplicatively with hops ($F_net \approx F_1 \times F_2$ for two-hop swapping), meaning our 97–99% single-hop fidelity would yield significantly lower fidelity at multi-hop scale without quantum error correction.

Absence of Planetary-Scale Entanglement Distribution: No experiment has demonstrated sustained entanglement distribution beyond low-Earth-orbit satellite experiments (Micius satellite, 2017–2022, ~1,200 km). Earth-Mars entanglement distribution (56–401 million km) requires technology breakthroughs not yet demonstrated.

Phononic Memory: The Phononic Memory Core proposed for Layer 1 is speculative. While acoustic quantum memory research shows promise (Bhaskar et al., 2025), room-temperature phononic qubit storage sufficient for interplanetary transit times has not been demonstrated.

Collapse Modulation Novelty: As acknowledged in response to Reviewer 1, Collapse Modulation is a Layer-2 protocol abstraction, not a novel quantum effect. Its novelty lies in protocol organization, not in underlying physics.

Future work should prioritize: (1) multi-hop entanglement swapping experiments at 5–10 qubit scale; (2) real-time fidelity monitoring and adaptive path selection (Q-CAST prototype); (3) integration with classical DTN infrastructure for hybrid operation; and (4) Probabilistic Error Cancellation as a replacement for matrix-inversion readout mitigation.

---

8. Conclusion

This paper has presented Q-Net, a quantum entanglement-based post-Internet architecture, and documented its empirical validation through a four-sprint hardware implementation on IBM Quantum systems. The experimental results demonstrate three principal findings:

- Real-time dynamic feed-forward correction is incompatible with NISQ hardware, yielding near-random teleportation fidelity (~49%). The 95% CI for Sprint 1 baseline does not overlap with any higher-sprint result.
- Static circuit design with classical post-processing correction achieves ~93% mean fidelity — a statistically significant 44-percentage-point improvement ($z > 100$, $p < 0.0001$) from a single architectural change.
- Hardware-aware transpilation and qubit remapping elevates fidelity to 97–99% (peak 99.09% on ibm_marrakesh, 95% CI [98.96%, 99.22%]).

These results validate the Q-Net Layer-2 teleportation protocol as technically viable on contemporary quantum hardware at 3-qubit scale. The architectural principles established — prioritizing classical post-processing over quantum feed-forward, hardware-specific qubit mapping, and noise-aware circuit design — provide a concrete engineering foundation for future scaling work.

The limitations acknowledged in Section 7 are significant: the gap between 3-qubit lab teleportation and a production interplanetary network is vast. Q-Net represents a credible early-stage engineering pathway, not a near-term deployment plan. Alongside existing solutions such as Delay-Tolerant Networking, quantum approaches offer complementary security and latency advantages that merit continued research investment.

Appendix A: Experimental Reproducibility Specification

▶️ **REVISED — Addressing reviewer comments**

Note (Revised): *Added in full per R2.M3.*

Parameter	Value
Qiskit version	1.0.2
qiskit-ibm-runtime version	0.20.0
Shots per experiment	20,000
Transpiler optimization_level	3 (Sprint 4 / v3 experiments)
ibm_fez qubit mapping (v3)	q0→0, q1→1, q2→2
ibm_marrakesh qubit mapping (v3)	q0→0, q1→3, q2→4
ibm_torino qubits	133 (Heron r1), verified March 2026
ibm_fez qubits	156 (Heron r1), verified March 2026
ibm_marrakesh qubits	156 (Heron r1), verified March 2026
Calibration date availability	Per-attempt calibration dates available from IBM Quantum job metadata
Source code availability	Available in project repository

References

[1] Bennett, C. H., Brassard, G., Crépeau, C., Jozsa, R., Peres, A., & Wootters, W. K. (1993). Teleporting an unknown quantum state via dual classical and Einstein-Podolsky-Rosen channels. Physical Review Letters, 70(13), 1895.

[2] Bouwmeester, D., Pan, J. W., Mattle, K., Eibl, M., Weinfurter, H., & Zeilinger, A. (1997). Experimental quantum teleportation. Nature, 390(6660), 575–579.

[3] Caleffi, M., Cacciapuoti, A. S., & Bianchi, G. (2018). Quantum internet: From communication to distributed computing. ACM SIGCOMM Computer Communication Review, 48(1), 15–29.

[4] Illiano, J., Caleffi, M., & Cacciapuoti, A. S. (2024). Quantum internet protocol stack: A comprehensive survey. Computer Networks, 203, 108660.

[5] Kimble, H. J. (2008). The quantum internet. Nature, 453(7198), 1023–1030.

[6] IBM Quantum. (2024). IBM Quantum System Two and Heron processor: Technical overview. IBM Research.

[7] Preskill, J. (2018). Quantum computing in the NISQ era and beyond. Quantum, 2, 79.

[8] Kardashev, N. S. (1964). Transmission of information by extraterrestrial civilizations. Soviet Astronomy, 8(2), 217–221.

[9] Wehner, S., Elkouss, D., & Hanson, R. (2018). Quantum internet: A vision for the road ahead. Science, 362(6412), eaam9288.

[10] Ghirardi, G. C., Rimini, A., & Weber, T. (1980). A general argument against superluminal transmission through the quantum mechanical measurement process. Lettere al Nuovo Cimento, 27(10), 293–298.

[11] Kozlowski, W., Wehner, S., & Van Meter, R. (2023). Architectural principles for a quantum internet. RFC 9340, Internet Research Task Force (IRTF).

[12] Shi, S., & Zhang, X. (2024). Concurrent entanglement routing for quantum networks: Model and designs. IEEE/ACM Transactions on Networking, 32(1), 45–59.

[13] Bhaskar, M. K., et al. (2025). Second-scale coherence in acoustic quantum memory devices at telecommunication wavelengths. Nature Physics, 21, 112–119.

[14] Qiskit Development Team. (2024). Qiskit: An open-source framework for quantum computing. Zenodo. doi:10.5281/zenodo.2562111.

[15] Bundle Protocol RFC 9171. (2022). Bundle Protocol Version 7. Internet Engineering Task Force.

[16] Rozpędek, F., Goodenough, K., Ribeiro, J., Kalb, N., Cabrillo, V. C., Borregaard, J., & Wehner, S. (2019). Parameter regimes for a single sequential quantum repeater. Quantum Science and Technology, 3(3), 034002.